

fsophy: An Open-Source Simulation of an SDA-OCT-Class Free-Space Optical Link Modem

Hossein Khoshnevis

Abstract—Free-space optical links are becoming the default high-capacity interconnect where fibre cannot go, the crosslinks of proliferated low-Earth-orbit constellations being the leading example, and the Space Development Agency’s Optical Communications Terminal (OCT) standard defines the interoperable waveform: on-off-keyed light at 2.5 Gbps and 1550 nm, a fixed modem frame, and the 5G NR LDPC as payload FEC. The standard is public and defined to the bit, yet no open implementation of its physical layer has been available. This paper presents *fsophy*, a pure-numpy, fully vectorised baseband simulation of such a modem. The transmitter implements the standard’s framing, scrambler, header convolutional code and rate-matched LDPC bit for bit. The vacuum channel carries what space actually does to the signal: $1/Z^2$ spreading over thousands of kilometres, frame-by-frame fading from microradian pointing jitter, and a symbol clock pulled tens of ppm by Doppler. The receiver is an EDFA-preamplified square-law front end simulated at field level, followed by blind DSP: Gardner timing recovery, preamble frame sync, per-frame level tracking, Gamma-likelihood soft demapping matched to the skewed beat-noise statistics, and soft- or hard-input LDPC decoding. In its ideal configuration the receiver lands within about 0.2 dB of the exact chi-square theory of preamplified OOK, and its rate-1/2 FEC cliff within 0.9 dB of the channel’s soft-decision capacity; with every impairment on, the cliff sits at 11 received photons per line bit, about 13.5 dB below the standard’s $25 \mu\text{W}/\text{m}^2$ compliance floor at an 8 cm aperture. The package is small enough to read and fast enough to sweep, and it is released for others to complete.

Index Terms—Free-space optical communication, satellite communication, LDPC codes, clock recovery, open-source software, simulation.

I. INTRODUCTION

FREE-space optical links have moved from demonstration missions to production networks, and their clearest example is the inter-satellite mesh of proliferated constellations: hundreds of laser terminals are in orbit, thousands are contracted, and the Space Development Agency (SDA) made interoperability a procurement requirement by publishing the Optical Communications Terminal (OCT) standard [1]. The standard fixes the line format to the bit: OOK-NRZ or Manchester light at 1553.33 or 1536.61 nm, symbol rates from 312.5 Mbaud to 2.5 Gbaud, a modem frame with a 64-bit preamble and a convolutionally protected header, a frame scrambler, and, for the payload, the 5G NR LDPC code [2] at four rates between 0.5 and 0.85. It also fixes the physics floor: at least 2.5 W transmitted through a beam no narrower than $15 \mu\text{rad}$ FWHM, and a post-FEC BER of 10^{-6} or better at an irradiance of $25 \mu\text{W}/\text{m}^2$, quoted at 5500 km. CCSDS is standardising a similar non-coherent waveform [3], [4]. The documents are public, yet a researcher who wants to reproduce

the standardised link end to end, from encoder to decision, has had to work from the PDFs.

fsophy provides that baseband physical layer as readable, vectorised code, the free-space sibling of the author’s ofmphy package for the 800G coherent fibre interface: the same discipline of bit-exact standard structures, a receiver that earns its numbers blind, and validation anchored on textbook theory, the standard’s own floors, and published sensitivities. The direct-detection link swaps the coherent receiver’s carrier problems for a clock problem, and the package is built around that: after square-law detection there is no carrier, but the symbol clock arrives tens of ppm off because the terminals can move at kilometres per second relative to each other, as satellites do, so timing recovery takes the place that carrier recovery held in the fibre package. The contributions are:

- a complete, inspectable transmitter–channel–receiver chain for the OCT line format, in pure numpy;
- a bit-exact open implementation of the standard’s framing and FEC: preamble, scrambler, header CRC-16 and rate-1/6 convolutional code, and the NR LDPC (base graph 1, $Z = 384$) with the standard’s four rate-matching modes and both decoders;
- a vacuum channel and preamplified front end simulated at field level, so beat noise, pointing fades and clock Doppler appear with their real statistics; and
- a validation methodology anchored on three sides: the exact chi-square theory of preamplified OOK, the standard’s compliance floor, and published photon-efficiency results.

The library is meant to grow by contribution: Section VII lists the blocks it does not yet have, in the hope that readers adopt them.

II. RELATED WORK

Free-space optical simulation has mostly lived at the link-budget level, in tools that integrate margins over orbits, or in publication-specific models of single effects such as pointing jitter [5]. The components are all well studied: the 5G NR LDPC [6], and preamplified receiver theory settled since Humblet and Azizoglu [7] and the first amplifier-link analyses [8]. What has been missing is the assembled artefact: the standard’s exact frame, its exact code, a physical channel and a blind receiver in one place that runs. CCSDS’s non-coherent coding book, which adopts an LDPC of its own, is still at draft stage [4], leaving the SDA standard as today’s anchor. *fsophy* is not a network planner and not a terminal design tool; it is a readable executable form of the standardised link.

III. TRANSMITTER AND CHANNEL

Fig. 1 shows the chain. Each modem frame carries 8416 payload bits and their CRC-32, which together fill the LDPC information block exactly; the coded frame is prefixed by the 64-bit preamble `0x53225B1D0D73DF03` (MSB first) and a 960-bit coded header, and everything after the preamble is XORed with the standard's scrambler, the maximal-length sequence of $x^{15} + x^{14} + 1$ restarted each frame from a fixed seed. The scrambler does three jobs at once: it balances the optical duty cycle, guarantees the transition density the timing detector feeds on, and keeps the frame's structure from imprinting on the blind estimators. Frames follow each other with no gaps, so the receiver sees one continuous synchronous stream, which is what its clock recovery needs. Fig. 2 draws the frame and the code around it.

A. Header and its code

The header exists so that a receiver can learn the frame's identity and PL_RATE *before* it can decode the payload, so it is protected far below the payload threshold: 128 field bits and their CRC-16 (CCITT polynomial, zero-initialised) are zero-tailed to 160 bits and passed through a non-systematic rate-1/6, constraint-length-7 convolutional code with generators (octal) 175, 171, 151, 133, 127, 117, giving 960 transmitted bits. A convolutional code is the right tool at this size: 160 bits give an iterative code nothing to converge on, while Viterbi decodes the 64-state trellis exactly, and six-fold redundancy holds at raw error rates far beyond where the LDPC gives up. The simulation uses it the way a terminal does: decoded headers whose CRC passes anchor the frame count of the whole synchronous stream, so frame identity is earned, not assumed.

B. Payload LDPC

The payload FEC is the terrestrial 5G NR data-channel code adopted unchanged: base graph 1, lifting factor $Z = 384$, information block $K = 22Z = 8448$. The parity-check matrix is the 46×68 base graph expanded by circulants, $H_{ij} = P^{s_{ij}}$ with P the $Z \times Z$ cyclic shift; encoding solves the four core parity blocks through the base graph's double-diagonal structure and reads the 42 extension parities directly off their rows. As in 5G, the first two systematic blocks (768 bits) are never transmitted, which is why the standard's rates come out as they do: PL_RATE 1 to 4 transmit the first 6, 9, 13 or 24 parity blocks (2304 to 9216 bits), for rates $8448/(7680 + \text{parity}) = 0.8462, 0.7586, 0.6667$ and 0.5000 , matching the standard's table exactly. The punctured blocks enter the decoder with zero LLR and are recovered by the iteration itself, which is also why plain bit-flipping is not usable on this code: it has nothing to start from on 768 bits. Two decoders are provided. The soft decoder is flooding normalised min-sum [6] with factor 0.75,

$$m_{c \rightarrow v} = 0.75 \left(\prod_{v' \in c \setminus v} \text{sgn } m_{v' \rightarrow c} \right) \min_{v' \in c \setminus v} |m_{v' \rightarrow c}|, \quad (1)$$

run on the sub-graph a PL_RATE actually pays for (rows whose parity never left the transmitter carry no information

and are dropped), with syndrome-based early termination. The 0.75 attenuation exists because the min approximation overestimates every check magnitude; scaling the messages down is the one-parameter fix, and it buys a second robustness that matters here, forgiveness of imperfect LLRs. The exact sum-product tanh rule ships as an option and was measured to gain nothing on this channel, losing slightly in the cliff region for exactly that reason: the Gamma LLRs are approximate, and exact belief propagation trusts them more than they deserve. The hard decoder is the same iteration fed only the sign of every channel bit; the gap between the two is then exactly what calibrated soft demapping is worth, the role iterated bounded-distance decoding played against Chase-Pyndiah in the fibre package.

C. Vacuum channel

Between the telescopes there is nothing, and the model says so: no dispersion, no turbulence, no polarisation drift. What remains is geometry, pointing and motion. The transmitted Gaussian beam of $1/e^2$ half-angle divergence θ_w (the standard quotes FWHM; $\theta_w = \theta_{\text{FWHM}}/\sqrt{2 \ln 2}$) spreads to a far-field radius $w = \theta_w Z$, giving on-axis irradiance

$$I(Z) = \frac{2P_{\text{tx}}\eta_{\text{tx}}}{\pi \theta_w^2 Z^2}, \quad \bar{n} = \frac{I(Z) A_{\text{rx}}\eta_{\text{rx}}}{h\nu R_b}, \quad (2)$$

with $\bar{n}$ the received photons per bit, the currency every result in this paper is quoted in. At the reference numbers (2.5 W, 15 μrad , 8 cm aperture, 5500 km) the budget delivers about 230 $\mu\text{W}/\text{m}^2$ and 2300 photons per bit, comfortably above the standard's 25 $\mu\text{W}/\text{m}^2$ floor: a floor, so that real terminals can spend the difference on losses and margin.

Pointing error moves the receiver off the beam axis, and the irradiance follows the Gaussian profile: a radial error θ_e costs a power factor $h = \exp(-2\theta_e^2/\theta_w^2)$. Each axis of the error is a Gauss-Markov process with microradian rms and a few hundred hertz of bandwidth, the residual of the fine-steering loop. Against microsecond frames that error is frozen: fading arrives frame by frame, not symbol by symbol, and with Rayleigh-distributed radial error the mean gain has the closed form $\mathbb{E}[h] = \gamma^2/(\gamma^2 + 1)$ with $\gamma = \theta_w/2\sigma_j$, which the fade-statistics unit test pins the simulation against. The whole fade distribution is just as closed, $P(h < x) = x^{\gamma^2}$, so the frame-outage floor of a faded link is one power law, $(\bar{n}_{\text{cliff}}/\bar{n})^{\gamma^2}$, drawn against measurement in Fig. 6. Relative motion, up to ± 7.5 km/s between crossing satellites, does two things: the optical Doppler shift, which square-law detection never sees, and the symbol-clock shift $\delta = v_r/c$ of up to ± 25 ppm, which the receiver very much sees and must track.

D. Electro-optic front ends

The transmitter has a finite extinction ratio, because a real modulator never fully closes: mark and space powers are $P_{1,0} = 2\bar{P} r^{\pm 1/2}/(r^{1/2} + r^{-1/2})$ for ratio r , 30 dB by default, plus laser relative intensity noise and a 4th-order Bessel bandwidth of $0.75R_b$. The receiver is an EDFA-preamplified square-law detector simulated at field level, and the preamplifier is what makes the receiver good: with 35 dB




Fig. 1. System overview: transmit path, the vacuum channel with the electro-optic front ends, and the blind receiver with its decoders. Timing recovery holds the place that carrier recovery held in the coherent fibre chain.




Fig. 2. The OCT modem frame at PL_RATE 4 (to scale), the BG1 codeword it is cut from, and the four standard rate-matching modes. The first two systematic blocks are never transmitted; the decoder recovers them.

of optical gain ahead of the photodiode, electronics noise stops mattering and sensitivity is set by the physics of the ASE instead. The signal field rides in one polarisation, amplified spontaneous emission rides in both with input-referred spectral density n_{sp} photons/s/Hz per polarisation ($n_{sp} = 10^{NF/10}/2$), and the photocurrent is

$$i(t) \propto |E_s(t) + n_x(t)|^2 + |n_y(t)|^2, \quad (3)$$

so signal-ASE and ASE-ASE beat noise appear with their real statistics rather than as bolted-on Gaussians. An optical filter of two-sided width B_o ahead of the detector admits $M = B_o T_b$

independent field samples per bit and polarisation, and each is a noise mode the chi-square statistics count ($M = 8$ at the defaults). After an electrical Bessel filter, an 8-bit ADC samples at two samples per symbol on its own clock, offset by the Doppler-and-oscillator ppm and by an arbitrary phase, which is where the receiver's real work begins. Two samples per symbol is a deliberate economy, priced into Table I: the interpolator works on a barely-Nyquist waveform, part of the front ends' 1.3 dB.

IV. RECEIVER DSP

The receiver is blind: frame identity comes from decoded headers, levels from the data itself, and timing from the waveform. The blocks run in the order that makes each one's assumptions true, and everything is vectorised over blocks of symbols rather than looped per symbol.

A. Timing recovery

The ADC stream is resampled onto the transmit symbol grid by a Gardner loop [9] driving a cubic interpolator [10]. Per block of 64 symbols the loop interpolates on-time and mid samples at its current phase and rate estimate, averages the timing error

$$e = \langle y_{k-1/2} (y_k - y_{k-1}) \rangle_{\text{block}}, \quad (4)$$

and updates a proportional-integral pair: the integrator absorbs the ppm-scale rate offset, the proportional path the residual phase. The mid sample is the whole trick: an NRZ waveform keeps its timing information in its transitions, and the error is odd around the eye centre, crossing zero exactly when the on-time samples sit where they should. The Gardner detector needs no decisions and no carrier, which is why it can run first; the scrambler upstream guarantees it transitions to work on. Measured pull-in covers ± 40 ppm, beyond the ± 25 ppm a LEO crosslink needs, and the loop stays locked down to the FEC cliff.

B. Frame synchronisation and identity

Correlating the (bipolar) 64-bit preamble over a two-frame window finds the frame boundary. Sixty-four bits are enough because correlation stacks the preamble's whole energy against single-sample noise, an 18 dB advantage over deciding any one bit. Decoded headers whose CRC-16 passes then anchor the frame count, and because the stream is synchronous every other frame follows by position. This is the receiver's whole alignment machinery, and it is the standard's: no genie index is used anywhere.

C. Level tracking and the little equalizer

OOK decisions need the two levels, and with pointing fades the levels breathe frame by frame, so each frame estimates its own: two-means clustering gives (μ_0, μ_1) and the per-level spreads (σ_0, σ_1) , which differ, because signal-ASE beating makes ones noisier than zeros. An optional 7-tap decision-directed feed-forward equalizer, block-LMS like every adaptive loop in the package, cleans up the residual intersymbol interference of the two Bessel front ends.

D. Gamma-likelihood demapping

The demapper turns each sample into the LLR the decoder deserves. Beat-noise densities are chi-square, right-skewed with level-dependent variance, and modelling each level as a Gamma density matched to its two moments,

$$L(u) = \log \frac{p_\Gamma(u; k_0, \theta_0)}{p_\Gamma(u; k_1, \theta_1)}, \quad k_i = \frac{\mu_i^2}{\sigma_i^2}, \quad \theta_i = \frac{\sigma_i^2}{\mu_i}, \quad (5)$$

places the decision crossing within counting noise of the sweep-optimal threshold. The Gaussian shortcut is kept as a fallback and as a lesson: its crossing sits visibly too high on the skewed densities, a measured factor three in pre-FEC BER at the sensitivity limit.

E. Decoding and counting

After descrambling, header Viterbi decoding and payload LDPC decoding, the per-frame CRC-32 delivers the standard's own good-frame verdict. Pre-FEC BER is counted on the transmitted coded bits, post-FEC on the 8416 payload bits per frame; the first frame of every run is acquisition, settling the timing loop and the level estimators, and is excluded from every count, which is how a real terminal behaves during link stand-up.

V. VALIDATION AND RESULTS

Validation is anchored three ways. First, theory: the reference receiver for preamplified OOK is integrate-and-dump behind an M -mode optical filter, and its decision statistic is exactly chi-square: on a one, a noncentral mode carrying the signal plus $2M-1$ central ASE modes; on a zero, the same with the extinction-ratio leakage as the noncentrality [7]. The package builds those densities numerically, reads the minimum-error threshold off them, and holds them as the analytic curve in Fig. 3; the same densities give the channel's soft-decision mutual information, so every code rate has its




Fig. 3. Measured sensitivity waterfall at PL_RATE 4 with every impairment on. Open markers are counting floors where no errors were observed. Pointing fades smear the soft cliff into a low-rate outage tail before the floor is reached.

capacity threshold in photons to hold the decoders against. Configured ideally (rectangular pulses, integrate-and-dump, 40 dB extinction, no RIN, full-rate ADC), the full DSP chain lands within about 0.2 dB of this exact curve, which pins the calibration of the field-level front end, the timing loop and the demapper at once.

Second, the standard. Fig. 3 shows the measured waterfall against received photons per line bit with every impairment on. With quiet pointing the soft rate-1/2 cliff sits at 11 photons per line bit, -54.5 dBm at 2.5 Gbps, which is 23 photons per information bit at the frame's 47% payload fraction; 2 μ rad of jitter stretches a thin outage tail out to about 17: rare frames caught by a deep fade, at rates that fall by an order of magnitude over the tail before the measured floors begin. The hard-input decoder needs roughly 5 photons more, and the pre-FEC curve runs about 1.7 dB from the exact theory of the ideal receiver. Table I decomposes that distance at a fixed pre-FEC operating point, and the decomposition is telling: the blind DSP itself costs almost nothing (0.04 dB on the ideal front end), the electro-optics cost 1.3 dB, and the channel's pointing loss the rest, so the penalty lives where a hardware designer would expect it, not in the algorithms. Translated through an 8 cm aperture, the standard's 25 μ W/m² floor delivers about 250 photons per bit, some 13 dB above the measured requirement; the floor is generous by design, since it must absorb every vendor's losses, ageing and pointing budget.

Two more measurements bound the design. The frame machinery is not the limit: headers and timing hold to about 7 photons per line bit, 2 dB beyond the payload cliff, so the code sets the sensitivity, as it should. And Manchester, measured the same way, pays about 2 dB against NRZ at the same chip rate (its cliff sits near 17 photons per information-carrying bit), the price of a guaranteed transition every bit.

Third, the physics around the numbers behaves. Fig. 4 folds the received waveform at the recovered timing: the ones are visibly noisier than the zeros, which is the chi-square asymmetry the Gamma demapper exists for. Sweeping

TABLE I
PHOTONS/BIT FOR PRE-FEC BER 10^{-3} (REGENERATED BY
DOCS/STANDARDS_GAP.PY).

Chain variant	photons/bit	Distance
Ideal receiver (chi-square theory)	29.7	—
Blind DSP on the ideal front end	29.8	+0.0 dB
+ ER, RIN, Bessel, 8-bit 2-sps ADC	40.1	+1.3 dB
Full chain (+ Doppler, pointing)	44.3	+1.7 dB
Standard floor at an 8 cm aperture	248	7.5 dB margin




Fig. 4. Received eyes folded at the timing the Gardner loop recovered: (a) at 400 photons/bit, where the asymmetric beat noise on the ones is plain; (b) at 16 photons/bit, near the soft cliff, where the decoders live.

range at the standard's transmit floors (Fig. 5), the reference terminal holds double-digit margin across every LEO crosslink geometry, and the line falls at exactly 20 dB per decade of range, the $1/Z^2$ law with nothing riding on top of it, because vacuum adds nothing else. It crosses the compliance point at 5500 km with about 23 dB in hand and does not run out of photons until beyond 60 000 km, which is why the same waveform class is credible for MEO and lunar-relay distances. Sweeping pointing jitter at fixed photons (Fig. 6) shows the fading mechanism: mean power falls slowly, but past $\sigma_j \approx \theta_w/4$ the frame-outage rate takes over long before the average BER looks alarming, which is exactly why the standard specifies its floor with margin rather than at the cliff. The dashed curve is the geometric floor $(\bar{n}_{\text{cliff}}/\bar{n})^2$; the measured outage rides above it, because a frame in a deep fade does not only lose photons, it drags the timing loop and the level estimators with it, and that excess is receiver dynamics, not geometry.

Against the literature, the numbers land where they should at every level, and this is checked, not asserted. The theory engine reproduces the textbook anchor: at one noise mode, a 3 dB noise figure and a polarisation filter it gives 38.4 photons per bit for an uncoded 10^{-9} , the classic quantum-limit figure of preamplified OOK [7], [8] (a unit test holds it there), and the 80 photons of the receiver modelled here is that number paying for its second polarisation, eight modes and 4.5 dB noise figure. On coding, the soft-decision mutual information of the channel puts the rate-1/2 capacity threshold at 7.2 photons per line bit. The measured cliff of the full blind DSP on the ideal front end is 8.8 photons, a 0.9 dB gap that is what normalised min-sum is known to give up on this code [6]; the default electro-optics add 1.0 dB (the cliff of 11), and everything beyond that is pointing fades.




Fig. 5. Link margin above the measured soft cliff against range, at the standard's transmit floors (2.5 W, 15 μ rad) and an 8 cm aperture. Crosses mark simulated post-FEC failures.




Fig. 6. Pointing jitter at 40 photons/bit: average pre-FEC BER degrades gently, but frame outages dominate once the jitter reaches a quarter of the beam's $1/e^2$ divergence. Open markers are outage counting floors.

Measured photon efficiencies bracket the result from both sides: 25 photons/bit DPSK is a landmark of the uncoded preamplified class [11], flight coherent BPSK terminals of the EDRS class operate in the tens [12], and the record efficiencies below 2 photons/bit belong to phase-sensitive amplification with strong coding [13]. A rate-1/2 OOK chain 0.9 dB from its own channel capacity sits squarely inside that band, and the soft-versus-hard gap of about 2 dB is what sign-only min-sum is expected to give up [6].

VI. IMPLEMENTATION NOTES

Everything is numpy; there is no compiled extension. The lifted LDPC graph vectorises naturally: an edge is one `np.roll`, and the $Z = 384$ parallel check equations of a base row process elementwise once every message is rolled into the check's index domain, batched over codewords. The adaptive loops (Gardner, LMS) update once per block on block-averaged errors, which leaves their fixed points unchanged and only narrows tracking bandwidth that a frame-quasi-static channel does not need. A 12-frame run of the full chain, both decoders included, takes well under a second on a laptop-class core, and the deep points of Fig. 3 carry up to 10^7 payload bits each. The repository ships 41 unit tests, including bit-level checks of the frame and code structures against the standard's

numbers and the chi-square theory match; every figure and table in this paper regenerates from committed scripts with fixed seeds.

VII. LIMITATIONS AND ROADMAP

The simulation begins where the beams already point at each other: the acquisition spiral, the 40/50 kHz tracking tone and the PAT state machine are not modelled, and neither is the standard's full header field allocation, of which the frame count, PL_RATE and encoding fields are implemented. Also absent: atmospheric turbulence and scintillation for links that dip into the atmosphere (the vacuum channel holds only above it), solar-conjunction background as more than a constant rate, avalanche-photodiode receivers for unamplified terminals, the standard's lower symbol rates (they run, but only 2.5 Gbaud is characterised), CCSDS O3K framing once its coding book leaves draft, and a coherent extension for the next generation of crosslinks, reuniting it with its fibre sibling. The code is structured so that each of these is one module or one config knob away.

VIII. CONCLUSION

fsophy makes the interoperable free-space optical physical layer executable: the OCT modem frame, header code and NR LDPC bit for bit, a vacuum channel with the impairments that actually matter in space, a field-level preamplified receiver, and blind DSP whose FEC cliff lands 0.9 dB from the channel's soft-decision capacity and 13.5 dB inside the standard's compliance floor. It is offered as a base that anyone can verify, break and extend.

REFERENCES

- [1] Space Development Agency, "Optical Communications Terminal (OCT) Standard," ver. 3.2.0, Apr. 2025. [Online]. Available: <https://www.sda.mil>
- [2] 3GPP, "NR; Multiplexing and channel coding," TS 38.212, Rel. 15, Dec. 2017.
- [3] CCSDS, "Optical Communications Physical Layer," Rec. 141.0-B-1, Aug. 2019.
- [4] CCSDS, "Non-Coherent Optical Communications Coding and Synchronization," draft Rec. 142.x, Feb. 2025.
- [5] M. Toyoshima, T. Jono, K. Nakagawa, and A. Yamamoto, "Optimum divergence angle of a Gaussian beam wave in the presence of random jitter in free-space laser communication systems," *J. Opt. Soc. Am. A*, vol. 19, no. 3, pp. 567–571, Mar. 2002.
- [6] J. Chen and M. P. C. Fossorier, "Near optimum universal belief propagation based decoding of low-density parity check codes," *IEEE Trans. Commun.*, vol. 50, no. 3, pp. 406–414, Mar. 2002.
- [7] P. A. Humblet and M. Azizoglu, "On the bit error rate of lightwave systems with optical amplifiers," *J. Lightw. Technol.*, vol. 9, no. 11, pp. 1576–1582, Nov. 1991.
- [8] N. A. Olsson, "Lightwave systems with optical amplifiers," *J. Lightw. Technol.*, vol. 7, no. 7, pp. 1071–1082, Jul. 1989.
- [9] F. M. Gardner, "A BPSK/QPSK timing-error detector for sampled receivers," *IEEE Trans. Commun.*, vol. 34, no. 5, pp. 423–429, May 1986.
- [10] C. W. Farrow, "A continuously variable digital delay element," in *Proc. IEEE Int. Symp. Circuits Syst. (ISCAS)*, Jun. 1988, pp. 2641–2645.
- [11] D. O. Caplan, M. L. Stevens, and J. J. Carney, "Demonstration of optical DPSK communication with 25 photons/bit receiver sensitivity," in *Proc. CLEO*, May 2006, paper CFH5.
- [12] F. Heine, H. Kämpfner, R. Meyer, and M. Lutzer, "Optical inter-satellite communication operational," in *Proc. IEEE MILCOM*, Oct. 2010, pp. 1583–1587.
- [13] R. Kakarla, J. Schröder, and P. A. Andrekson, "One photon-per-bit receiver using near-noiseless phase-sensitive amplification," *Light Sci. Appl.*, vol. 9, 153, Sep. 2020.